



COT 6930 - Generative Intelligence  
and Software Development Lifecycle

# Topic 1 - Generative Intelligence and Software Development Lifecycles

Dr. Fernando Koch  
kochf@fau.edu  
<http://www.fernandokoch.me>



---

# Agenda

- Introduction to GenAI
- Impact of GenAI on SDLC
- Use Cases
- Exercises





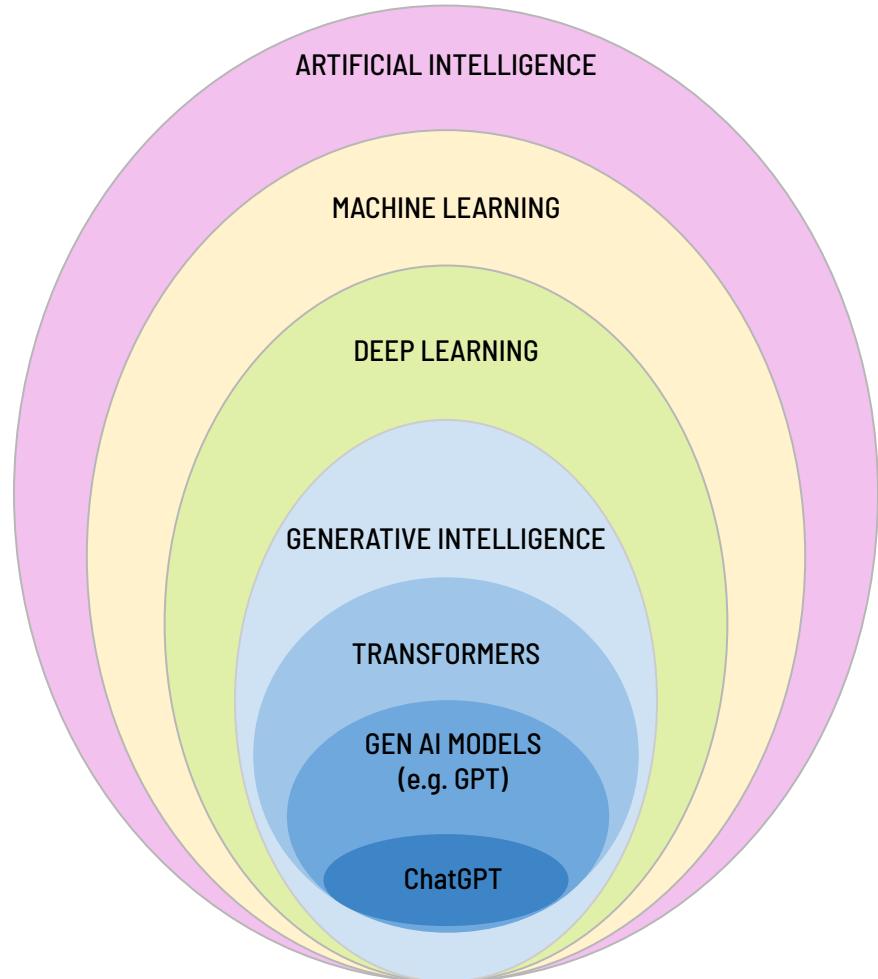
## Our Key Question

*How do Generative Intelligence Systems impact software development lifecycle by enhancing productivity and creativity across development stages?*





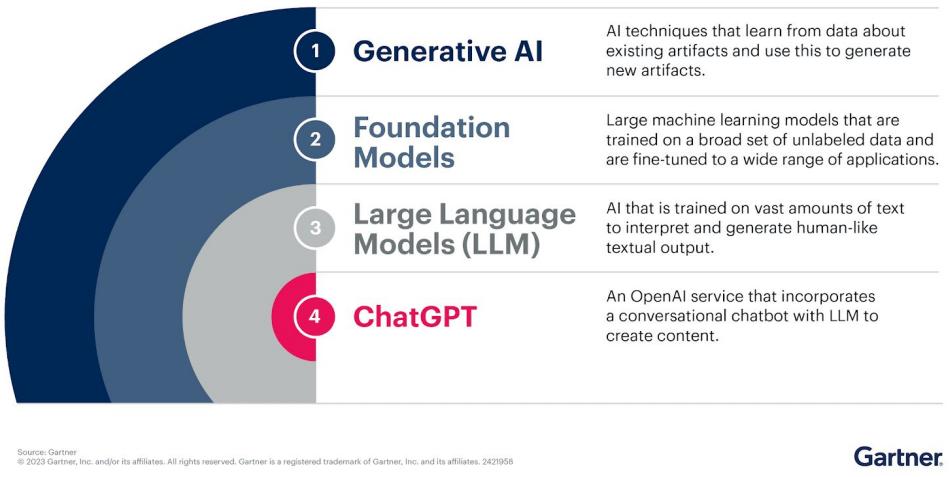
## *What is Generative AI?*



# Generative AI

## AI that creates new, coherent content

- Generates text, code, images, audio, and more
- Trained on large datasets to learn patterns and structures
- Foundation Models act as the core engines powering generative capabilities
- Capable of producing original outputs, not just 'retrieval'
- Applications span creative writing, design, coding, and synthesis
- Examples: GPT, DALL·E, Stable Diffusion



Gartner

# Why is Generative AI different?

1950s-

## Computation Models

If-Then-Else  
Finite State Machines  
Hardcode Rules

1970s-

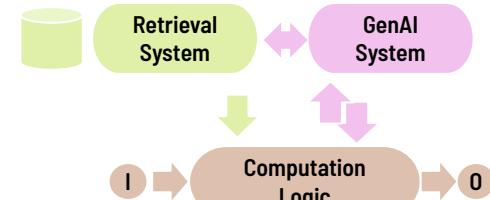
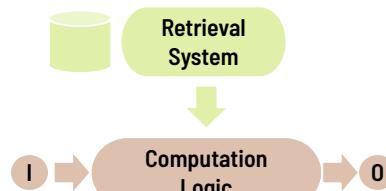
## Retrieval Models

Databases  
Rule-Based Systems  
Decision Trees

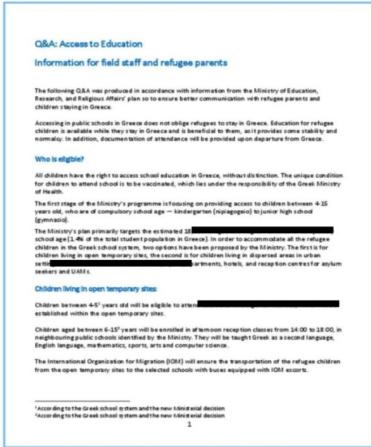
2010-

## Generative Models

Large Language Models  
Generative Adversarial Networks  
Diffusion Models

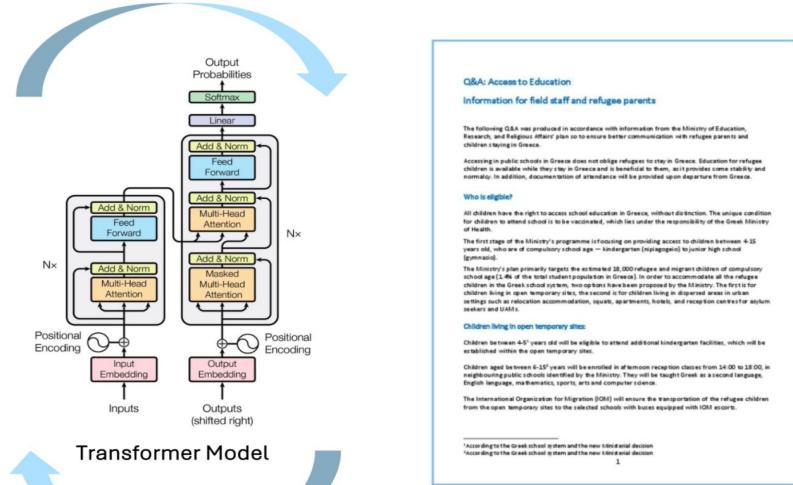


# The Process of Creating Generative Models



(A) Document with Obfuscated Parts

**(1) Inference:**  
Generate the content to complete the document based on predictions by the Model.



(B) Completed Document

**(2) Training:**  
Learn how to complete the missing parts provided enough examples.

## Traditional AI

- Focus on classification, prediction, and decision rules
- Relies on structured data (e.g., tables, labels, metrics)
- Uses supervised learning with predefined outcomes

## Generative AI

- Creates novel data like images, text, and code
- Learns from unstructured sources using self-supervised techniques
- Often unsupervised or self-supervised

# What Generative AI can do?

## The most common GenAI tasks today

Retrieval-augmented generation	<p>Based on documents or dynamic content, create a chatbot or question-answering feature.</p> <p>Building a Q&amp;A resource from a broad knowledge base, providing customer service assistance</p>	<p>Summarization</p> <p>Transform text with domain-specific content into personalized overviews that capture key points.</p> <p>Conversation summaries, insurance coverage, meeting transcripts, contract information</p>	<p>Content generation</p> <p>Generate text content for a specific purpose.</p> <p>Source Code, SQL, descriptions, posts, articles, email, others</p>
Named entity recognition	<p>Identify and extract essential information from unstructured text.</p> <p>Audit acceleration, SEC 10K fact extraction</p>	<p>Insight extraction</p> <p>Analyze existing unstructured text content to surface insights in specialized domain areas.</p> <p>Medical diagnosis support, user research findings</p>	<p>Classification</p> <p>Read and classify written input with as few as zero examples.</p> <p>Sorting of customer complaints, threat and vulnerability classification, sentiment analysis, customer segmentation</p>



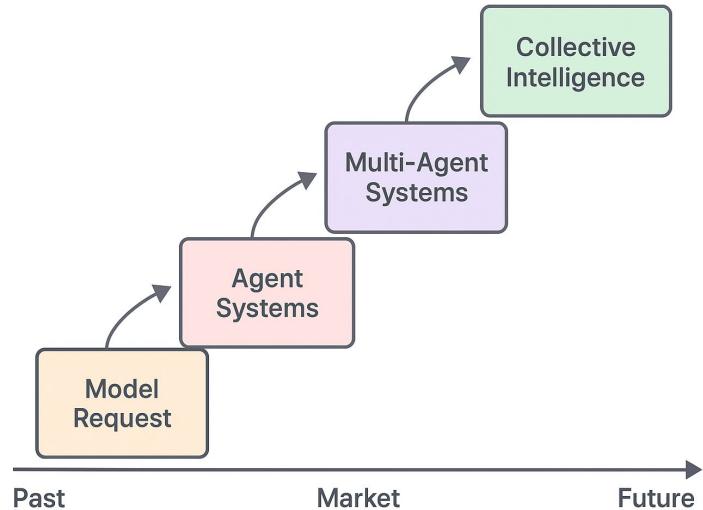
## Where GenAI does not help?

# When Generative AI Is and Is Not Effective

Use-case family	Generative models' current usefulness	Example use cases
Prediction/forecasting	Low	Risk prediction, customer churn prediction, sales/demand forecasting
Decision intelligence	Low	Decision support, augmentation, automation
Segmentation/classification	Medium	Clustering, customer segmentation, object classification
Recommendation systems	Medium	Recommendation engine, personalized advice, next best action
Content generation	High	Text generation, image and video generation, synthetic data
Conversational user interfaces	High	Virtual assistant, chatbot, digital worker

Source: Gartner  
© 2024 Gartner, Inc. and/or its affiliates. All rights reserved. 2937191

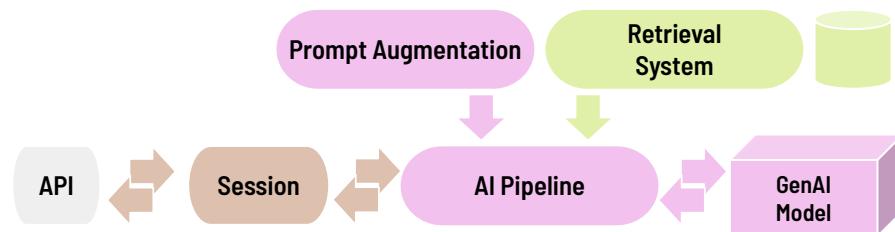
# Evolution of Utilization of Generative AI



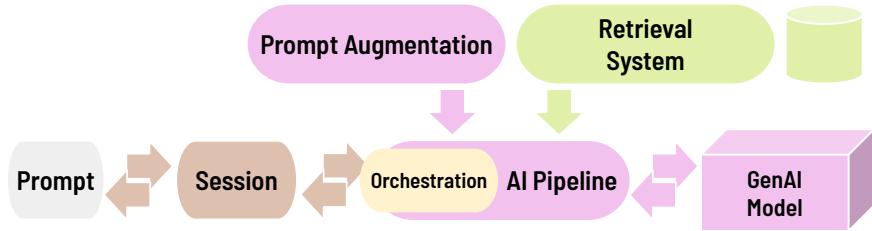
# Generative Intelligence System

End-to-end systems with autonomous content generation

- Session-aware workflows.
- Automated prompt engineering.
- Context-aware content generation
- Adaptive behaviour combining orchestration, memory, retrieval augmentation, others.
- End-to-end automation for LGM consumption.



# Generative Intelligence System



```
# ----- API handler -----
repo = SessionRepository()

def api_handler(input_message: str, context: str) -> str:
    # 0) Load session by context
    session = repo.get(context)

    # 1) Accept new request into session
    session.current.request = input_message

    # 2) Prompt Augmentation (build the ◇ prompt)
    pre = "You are a helpful assistant.\n== BEGIN USER REQUEST ==\n"
    post = "\n== END USER REQUEST ==\nProvide a clear, concise reply."
    session.current.prompt = pre + session.current.request + post    # ◇ highlight

    # 3) Retrieval (optional) → enrich the ◇ prompt
    ctx = retrieval_system(query=session.current.request)
    if ctx:
        session.current.prompt += f"\n\nContext:\n{ctx}"           # ◇ keep prompt

    # 4) AI Pipeline → GenAI model (configurable)
    llm = GenAIModel(
        model="llama3",          # or "llama3:70b"
        temperature=0.3,
        num_ctx=8192,            # acts like context window for Ollama
        num_predict=512
    )
    session.current.response = llm.generate(session.current.prompt)

    # 5) Persist session and return response to chatbot
    repo.save(session)
    return session.current.response

# ----- Example -----
reply = api_handler(
    input_message="What is LangGraph?",
    context="tenantA:user42:chat1" # the routing/index key
)
print(reply)
```

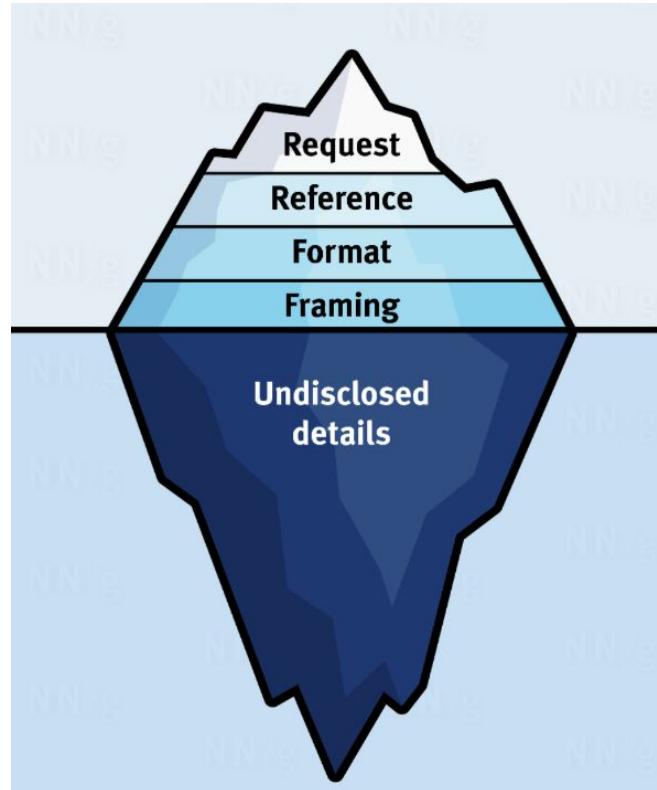


# What are Prompts?

Prompts are the inputs or instructions given to a generative AI model to guide its response.

- Act as the "query" or "command" for the AI
- Shape the style, tone, and depth of the output
- Can be questions, statements, or structured instructions
- The quality of a prompt directly affects the quality of the response

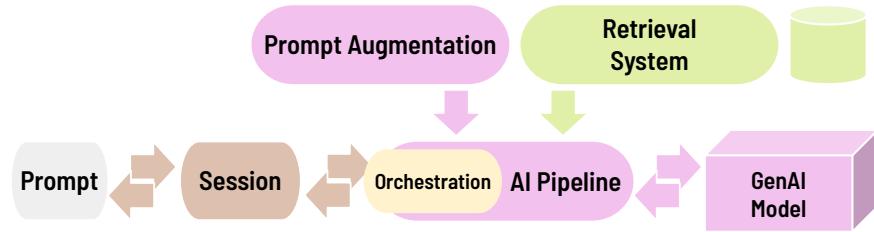
Note: we will have a deeper understanding of 'the math' behind Prompts as part of our deep dive in Topic 2: Large Generative Models.



# What are Sessions?

A session is the ongoing interaction between a user (or system) and the AI model.

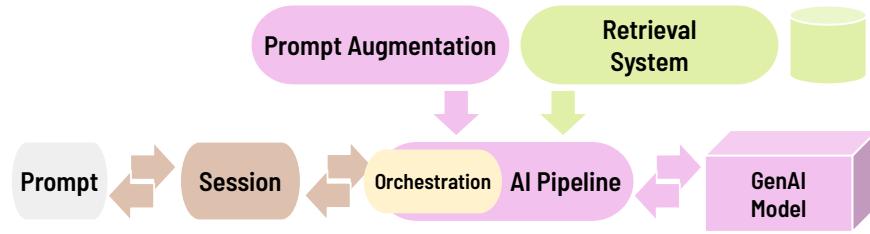
- Tracks the sequence of prompts and responses
- Maintains context across multiple turns
- Enables refinement and iteration over answers
- Critical for conversational AI and complex workflows



# What is an AI Pipeline?

Structured sequence that transforms user inputs into model responses.

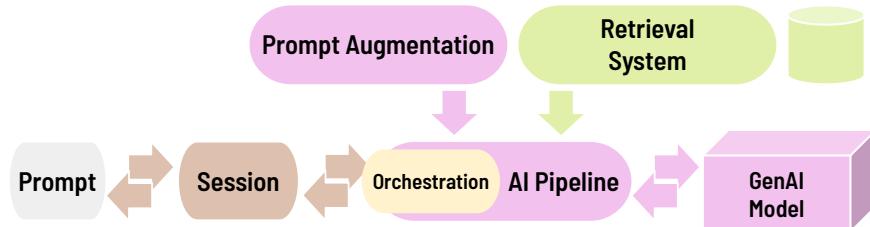
- Links prompts, retrieval, and model interaction
- Handles data preparation and enrichment
- Supports modular composition of components
- Enables repeatable and adaptive workflows



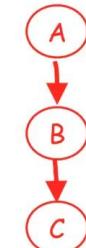
# What are Orchestration Workflows?

Frameworks to manage AI pipeline execution.

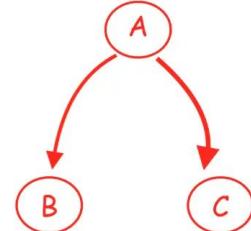
- Schedule, route, and monitor tasks
  - Enable conditional branching and looping
  - Support multi-step dependencies
- 
- **LangChain**: modular chaining of model calls
  - **LangGraph**: graph-based orchestration for adaptive flows



 **LangChain**



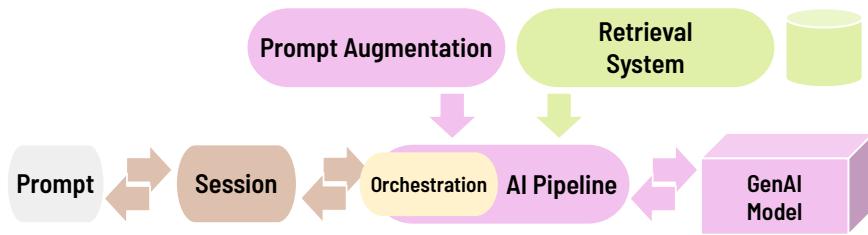
 **LangGraph**



# What's the role of the Retrieval System?

Mechanism to fetch domain-specific data for augmentation.

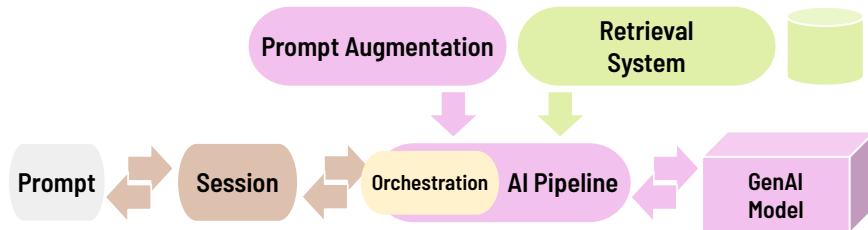
- Enhances prompts with relevant information
- Reduces model hallucinations
- Ensures context-aware responses
- Example:
  - a. Database retrieval
  - b. Knowledge-Based Retrieval
  - c. Knowledge Tree, others



# How to connect to GenAI Models?

Methods to access generative AI capabilities.

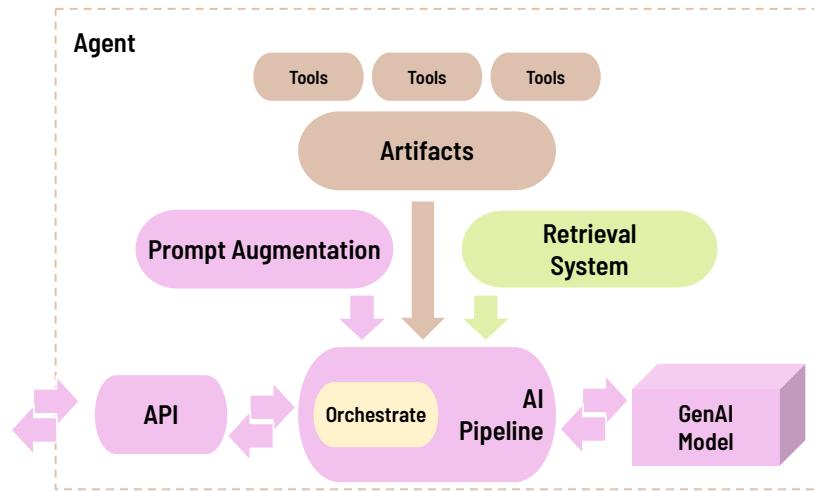
- Cloud API endpoints (e.g., OpenAI, Anthropic)
- Local API endpoints (e.g. Ollama)
- SDK endpoints
- Request parameters:
  - a. Context Window
  - b. max\_output\_tokens / max\_tokens.
  - c. Temperature
  - d. top\_p, top\_k, frequency/presence\_penalty



# What are Agentic Systems?

**Agentic systems =**  
**LLMs + tools + memory + autonomy.**

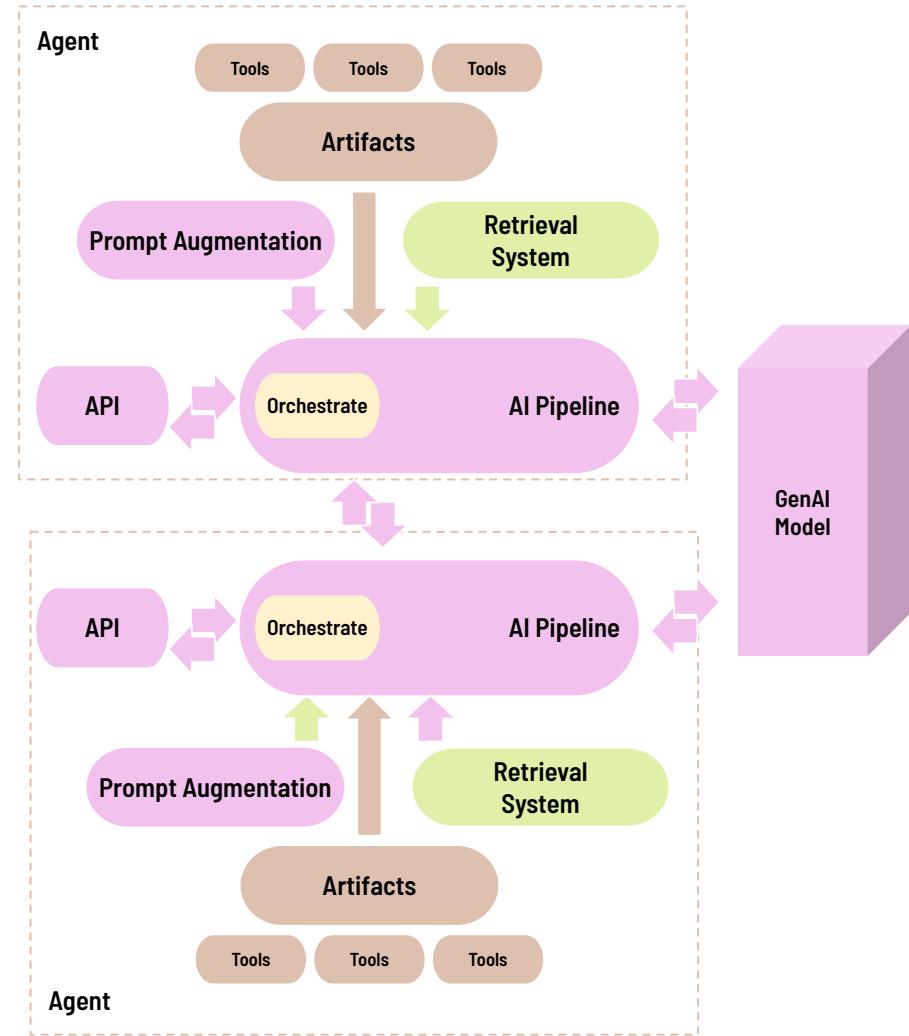
- **LLM-driven components capable of planning,** executing, and reasoning over sequences of actions.
- **Tools** provide external capabilities (e.g., APIs, search, code execution)
- **Prompt Augmentation:** Dynamically reformulates inputs based on past state, tools, and user goals.
- **Orchestration Layer:** Manages decisions, branching logic, and execution flows



# What are Multi-Agent Systems?

Composed of multiple autonomous agents that collaborate, delegate, or compete to accomplish complex tasks. Each agent:

- Has its own goal, memory, tools, and orchestration logic.
- May perform specialized roles (e.g., planner, coder, tester, executor).
- Can communicate with other agents via shared context or message passing.



---

# How people are using GenAI?



Latest Magazine Topics Podcasts Store The Big Idea Data & Visuals Case Selections



<https://hbr.org/2024/03/how-people-are-really-using-genai>

Generative AI

## How People Are Really Using Gen AI in 2025

by Marc Zao-Sanders

April 9, 2025



Illustration by Andrea Ucini

**Summary.** Last year, HBR published a piece on how people are using gen AI. Much has happened over the past 12 months. We now have Custom GPTs—AI tailored for narrower sets of requirements. New kids are on the block, such as DeepSeek and Grok, providing more... [more](#)

X  
Post

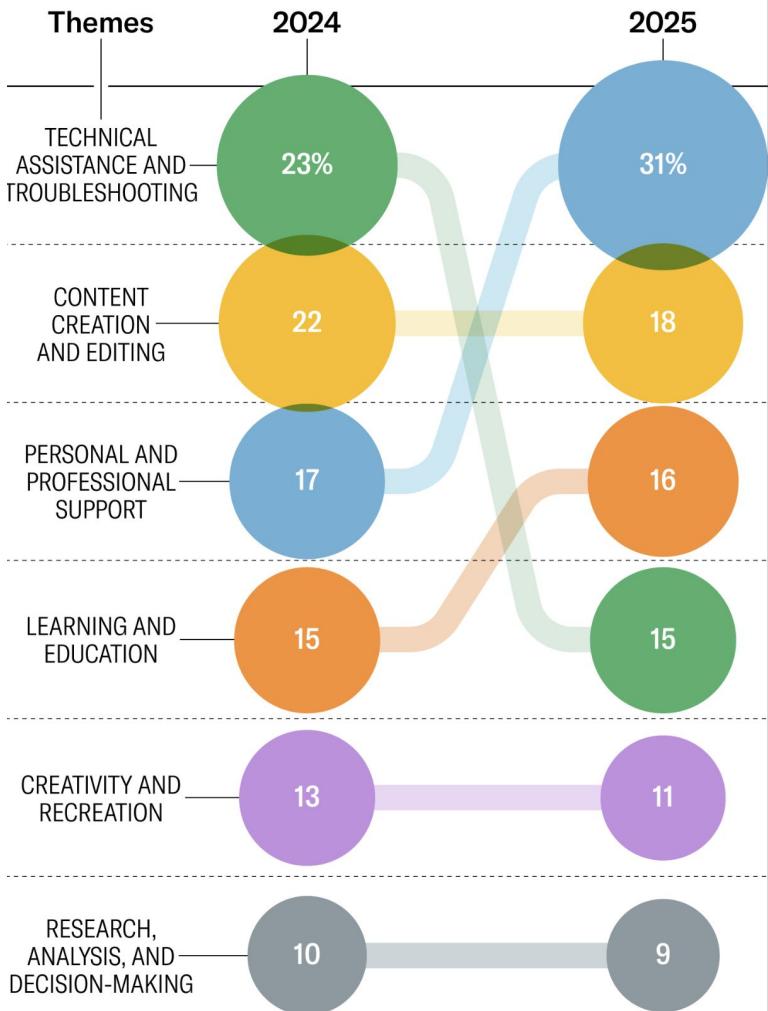
f  
Post

in  
Share

A year ago, I wrote a piece here about [how people were really using gen AI](#). That article seemed to hit a note: It was popular, featured in viral posts, and the [beautiful accompanying infographic](#) has been shared far and wide. The use cases split almost equally between personal and business needs, with roughly half spanning both.



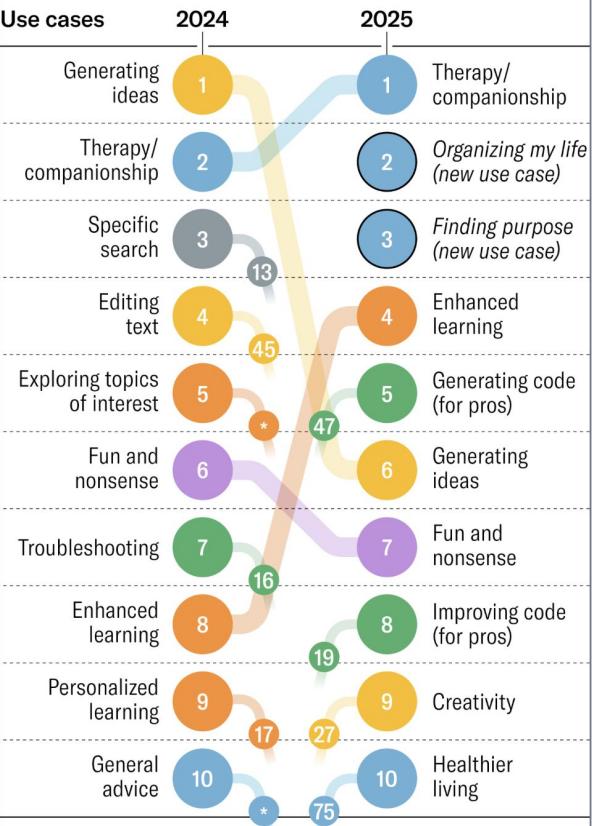
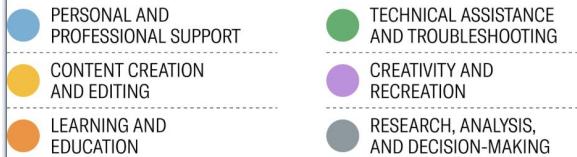
<https://hbr.org/2024/03/how-people-are-really-using-genai>





<https://hbr.org/2024/03/how-people-are-really-using-genai>

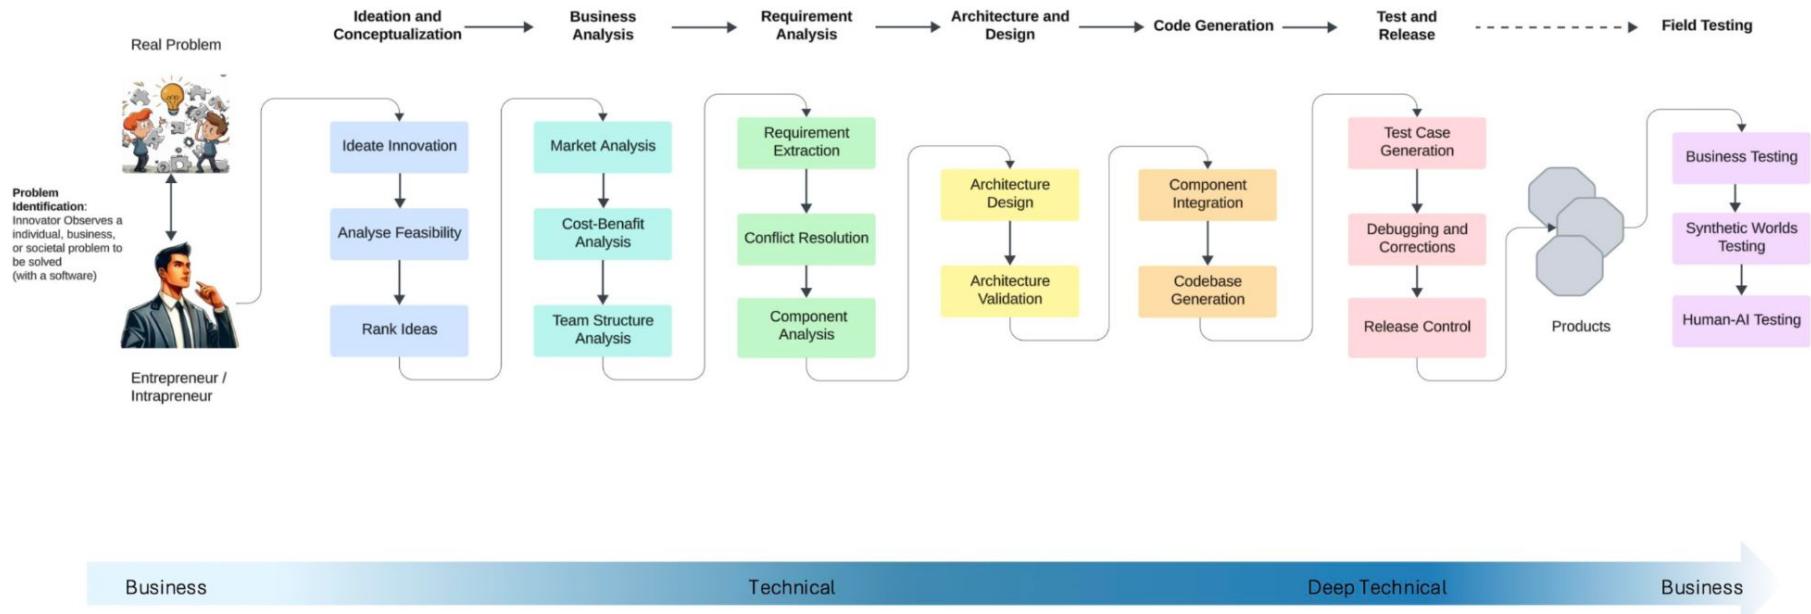
## Themes



---

# GenAI and SDLC

# Software Development Lifecycles



# Steps in the SDLC Process



## Problem Identification

- Understand stakeholder needs
- Define the core problem
- Align with organizational objectives

## Solution Ideation

- Collaborate to generate diverse ideas
- Evaluate feasibility and impact
- Develop a conceptual framework

## Requirement Analysis

- Clarify detailed requirements
- Prioritize requirements
- Define acceptance criteria

## System Design

- Translate requirements into system architecture.
- Incorporate technologies
- Ensure alignment between design and objectives

## Solution Development

- Adopt best practices for coding and architecture
- Build modular and scalable components
- Focus on seamless integration and testing



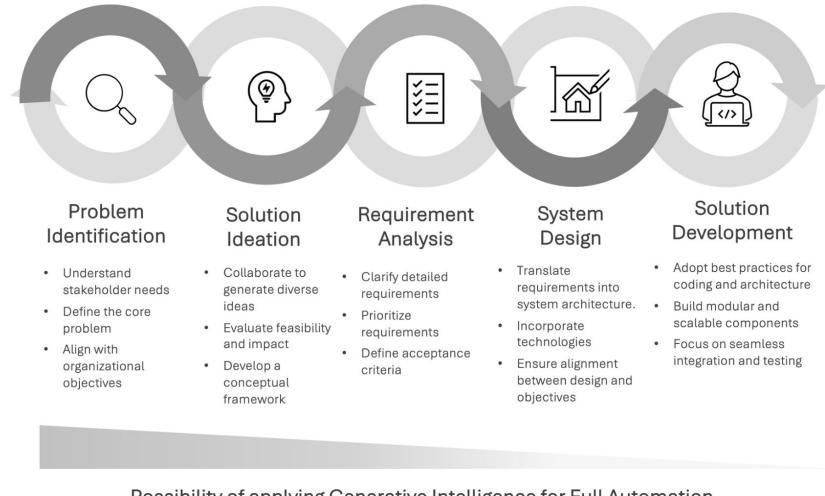
# Impact of GenAI on the SDLC Process

## 1. Problem Identification

- Analyze user feedback and logs to surface pain points.
- Cluster user needs using LLMs + embedding-based retrieval.

## 2. Solution Ideation

- Brainstorm multiple design ideas using prompt-based generation.
- Generate conceptual frameworks or architectural option



# Impact of GenAI on the SDLC Process

## 3. Requirement Analysis

- Convert informal business needs into formal requirements.
- Generate user stories and acceptance criteria.

## 4. System Design

- Propose system architectures aligned with requirements.
- Visualize component interactions or data flows using AI-generated diagrams.

## 5. Solution Development



---

# EXERCISE



---

## Exercise 1 - Getting Started with MyBot

**Objective:** Apply OwlMind Framework to develop a basic conversational bot to familiarize with GenAI interfaces and simple interactions.

<https://github.com/genilab/owlmind/wiki/Exercises>



## COT 6930 - Generative Intelligence and Software Development Lifecycles

**Dr. Fernando Koch**

kochf@fau.edu

<http://www.fernandokoch.me>