



COT 6930 - Generative Intelligence  
and Software Development Lifecycle

## Topic 5 - Generative Intelligence Pipelines

Dr. Fernando Koch  
kochf@fau.edu  
<http://www.fernandokoch.me>



---

# Agenda

Fundamentals of Generative AI Pipelines

Dynamic Orchestration

Adaptive AI Pipelines

Use Cases





## Our Key Question

*How can we design and manage AI pipelines that integrate multiple models, tools, and components to build intelligent, adaptive systems?*



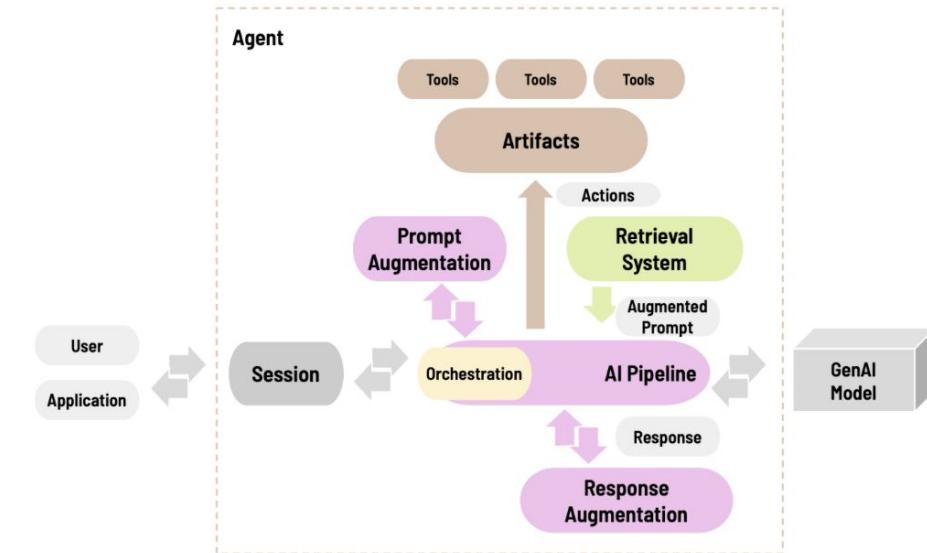
---

# Fundamentals of Generative AI Pipelines

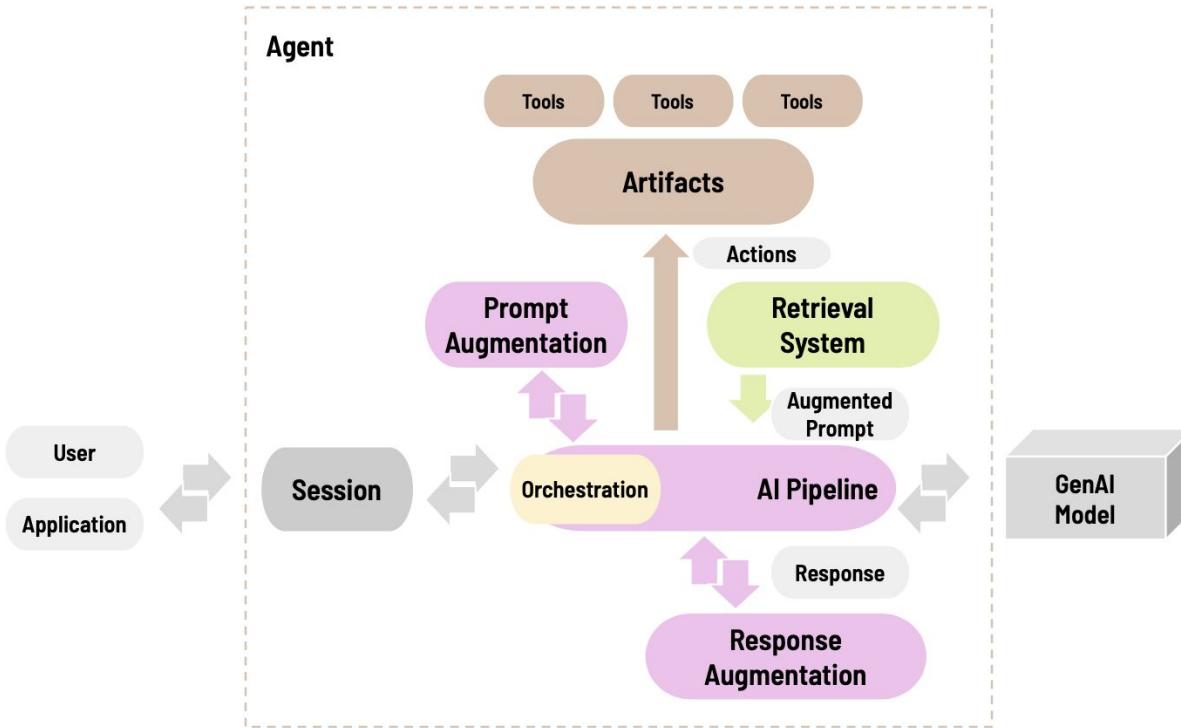
# Generative Intelligence System

End-to-end systems with autonomous content generation

- Session-aware workflows.
- Automated prompt engineering.
- Context-aware content generation
- Adaptive behaviour combining orchestration, memory, retrieval augmentation, others.
- End-to-end automation for LGM consumption.



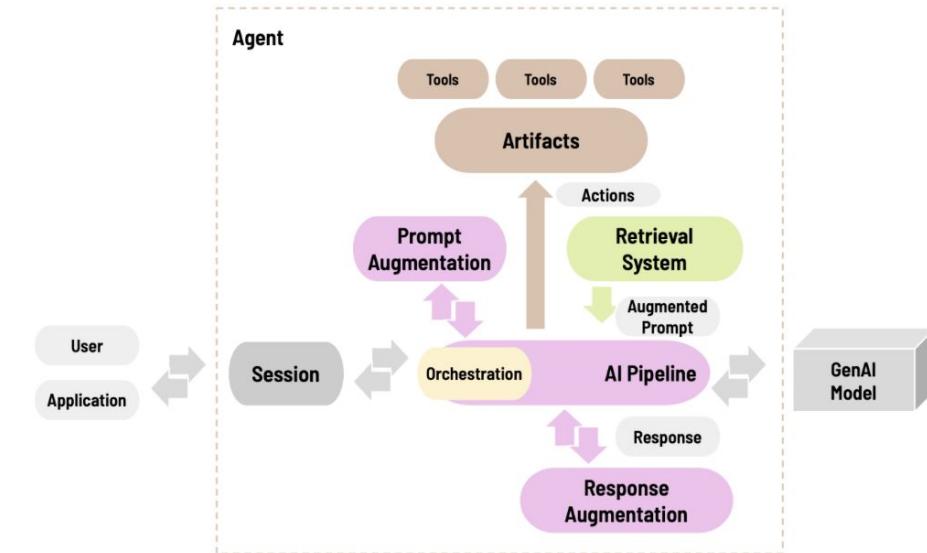
# Systematic view of GenAI Pipelines



# GenAI Pipeline

Structured workflow that connects AI models, tools, and knowledge sources to solve complex tasks.

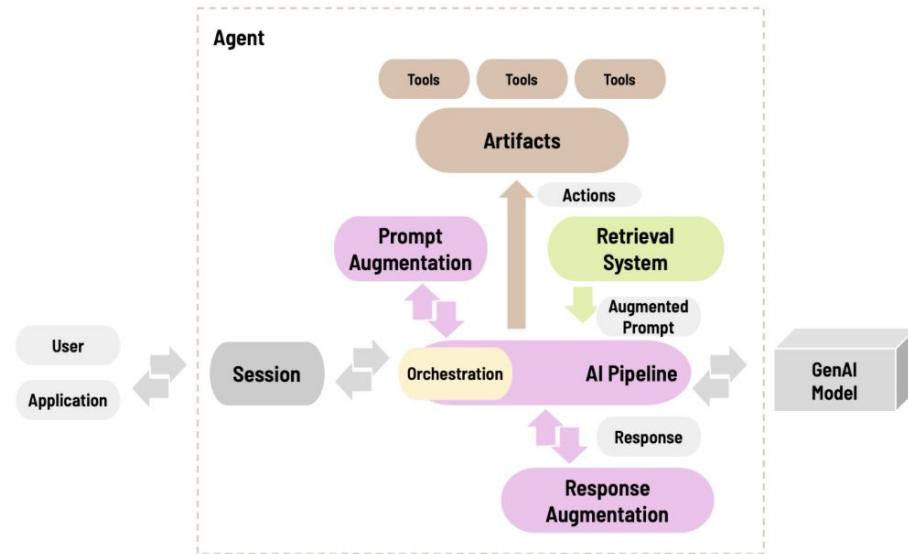
- Orchestrate processing of prompt requests through Session → Prompt Augmentation → Retrieval Systems → Artifacts → Model Requests.
- **Answer to 'What to do next?'**



# What are Sessions?

A session is the ongoing interaction between a user (or system) and the AI model.

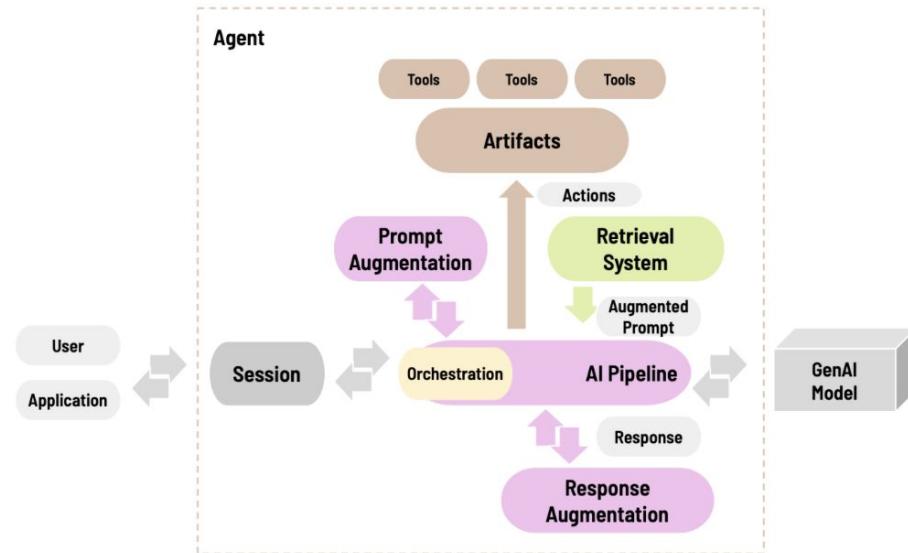
- Tracks the sequence of prompts and responses
- Maintains context across multiple turns
- Enables refinement and iteration over answers
- Critical for conversational AI and complex workflows



# What are the Orchestration Mechanisms?

Coordinates how different pipeline components interact to produce coherent results.

- **Answer to 'What to do next?'**
- **Flow Control:** Directs prompt requests through augmentation, retrieval, and response stages.
- **Adaptive Routing:** Dynamically chooses which systems (retrieval, tools, models) to invoke.
- **Error Handling & Validation:** Ensures robustness by refining or retrying when needed.



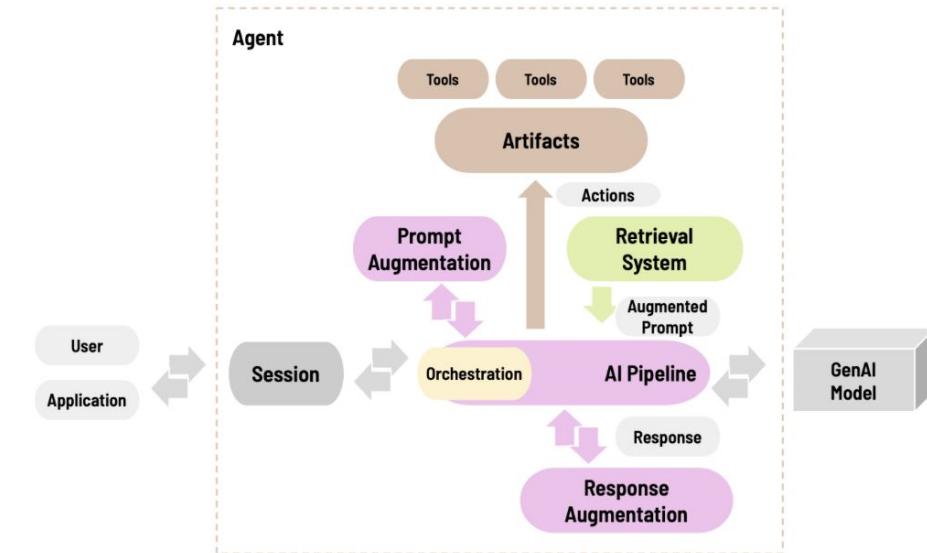
# Static vs Dynamic Orchestration

## Static Orchestration:

- Predefined, rule-based sequence of steps.
- Suitable for repetitive, predictable workflows.
- Example: Always → Retrieval → Model → Response Augmentation.

## Dynamic Orchestration:

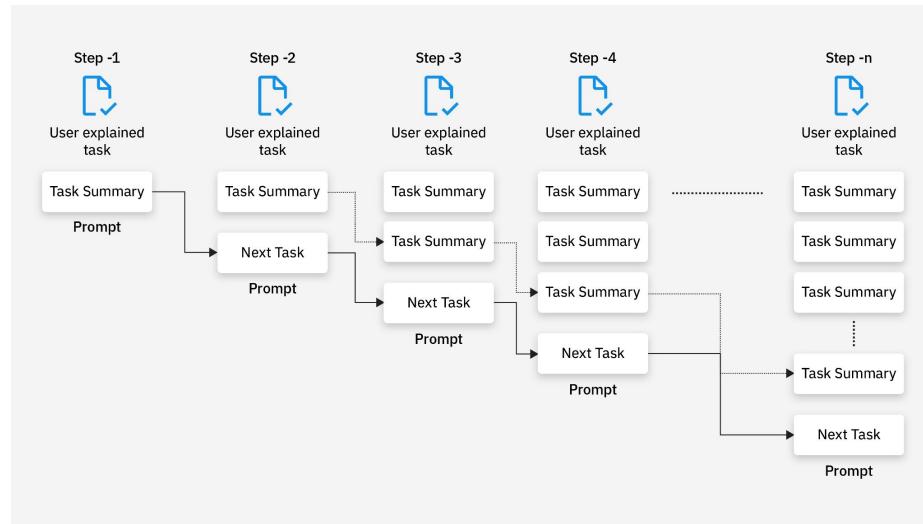
- Adaptive routing based on context, user input, or intermediate outputs.
- Enables flexibility and personalization.
- Example: Skip retrieval if session memory already has needed info.



# What is Prompt Chaining?

Breaks complex tasks into smaller, sequential steps where each output informs the next.

- **Stepwise Reasoning:** Output of one prompt feeds into the next for structured progress.
- **Task Decomposition:** Simplifies large problems into manageable subtasks.
- **Context Building:** Each chain stage enriches the prompt with intermediate results.



# LangChain and LangGraph

## LangChain:

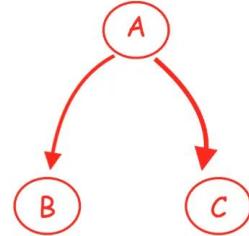
- Provides modular building blocks for chaining prompts, retrieval, tools, and models.
- Supports prompt templates, memory, and structured workflows.

<https://www.langchain.com/>

## LangGraph:

- Extends LangChain with graph-based orchestration of multi-step reasoning.
- Enables branching, looping, and conditional prompt flows.

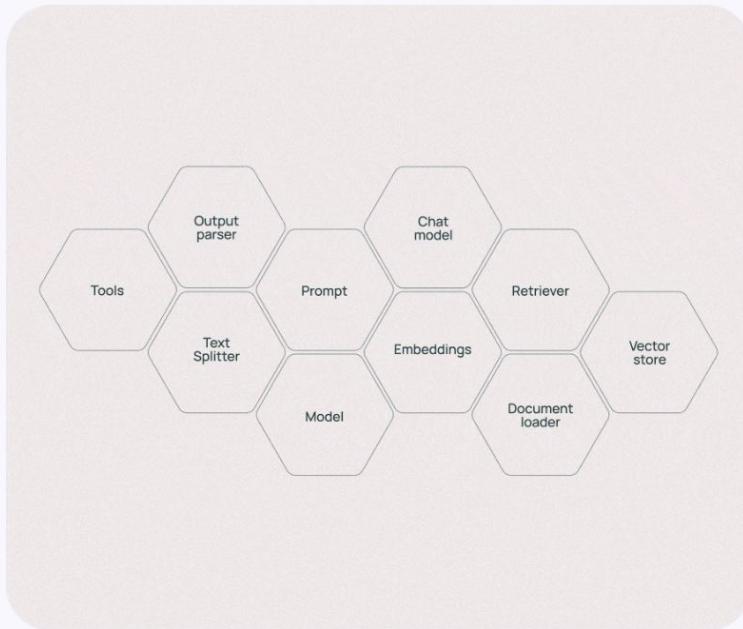
<https://www.langchain.com/langgraph>



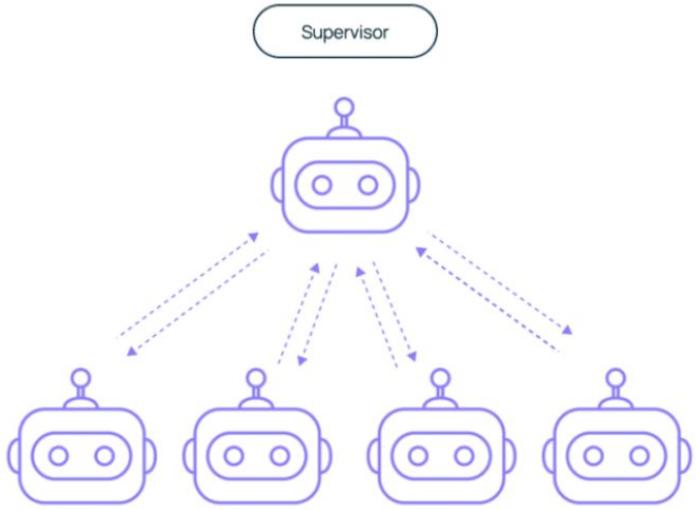
## Composable by design

LangChain's standard interface lets you experiment with different providers, tools, and databases – creating DevEx parity when gaps exist.

# Why use LangChain?



# LangGraph



## Build expressive, customizable agent workflows.

LangGraph's low-level primitives provide the flexibility needed to create fully customizable agents. Design diverse control flows – single, multi-agent, hierarchical – all using one framework.

See different agent architectures ↗

My Projects / Basic Prompting

Show Beta

Show Legacy

Search

I/O

{} Prompts

@ Models

Data

Processing

Vector Stores

Agents

Logic

\* Helpers

Calculator

Current Date

Message History

Bundles

Discover more components

+ New Custom Component

Playground Share

The diagram illustrates a workflow for generating a language model response. It starts with a "Chat Input" component (15ms) that receives a "Hello" message. This message is passed to a "Prompt" component (3ms), which contains a template: "Answer the user as if you were a GenAI expert, enthusiastic about helping them get started building something fresh." The output of the "Prompt" component is then fed into a "Language Model" component (2.14s). The "Language Model" component uses the "OpenAI" provider with the "gpt-4o-mini" model and an OpenAI API key. It also receives an "Input" message ("Receiving input") and a "System Message" ("Receiving input"). The final output of the "Language Model" is a "Model Response". This response is then displayed in a "Chat Output" component (14ms), which shows the generated response message.

Chat Input 15ms

Get chat inputs from the Playground.

Input Text

Hello

Chat Message

{} Prompt 3ms

Create a prompt template with dynamic variables.

Template

Answer the user as if you were a GenAI expert, enthusiastic about helping them get started building something fresh.

Prompt

Language Model 2.14s

Runs a language model given a specified provider.

Model Provider

OpenAI

Model Name

gpt-4o-mini

OpenAI API Key

Input

Receiving input

System Message

Receiving input

Model Response

Chat Output 14ms

Display a chat message in the Playground.

Inputs\*

Receiving input

Output Message

---

# LangFlow and Low-Code Solutions

Frameworks to manage AI pipeline execution.

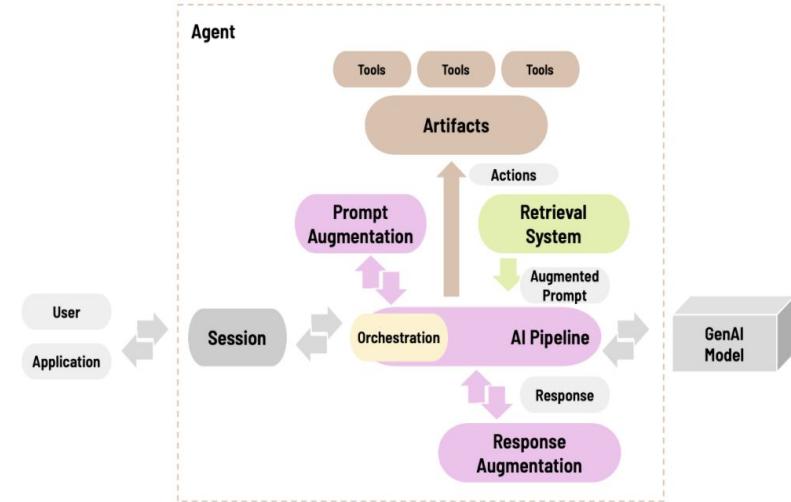
- Schedule, route, and monitor tasks
  - Enable conditional branching and looping
  - Support multi-step dependencies
- 
- **LangChain**: modular chaining of model calls
  - **LangGraph**: graph-based orchestration for adaptive flows

---

# Dynamic Orchestration

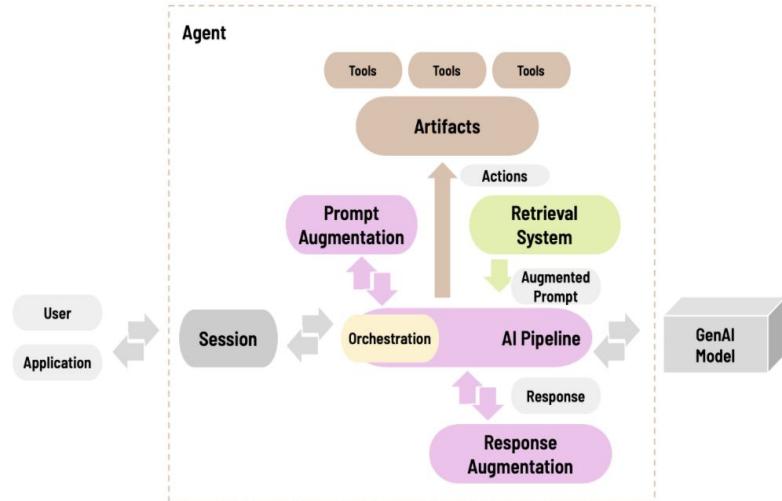
# Where can we implement Dynamic Orchestration?

- **Session Context** – Adjust orchestration based on session state
- **Prompt Augmentation** – Use different strategies per task
- **Retrieval Systems** – Adaptive retrieval depth/scope
- **Model Parameters** – Dynamic tuning of generation settings
- **Model Selection** – Use specialized models for different requests
- **Response Augmentation** – Adapt presentation style



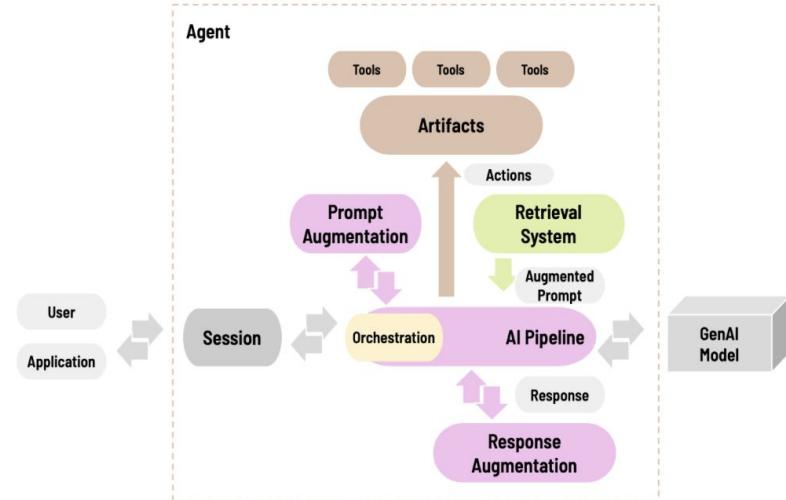
# Dynamic Orchestration of Session Control (1/2)

- **Topic Continuation** – If the user continues on the same subject, then reuse prior responses as context.
- **Topic Reset** – If a new topic is detected, then clear memory or down-weight old session data.
- **Adaptive Memory Length** – If the conversation is long or short, then expand or shrink the context window accordingly.
- **Relevance Filtering** – If past turns are not relevant to the current query, then exclude them from the prompt.
- **Context Summarization** – If the session history is too large, then summarize it into compact memory before injecting.



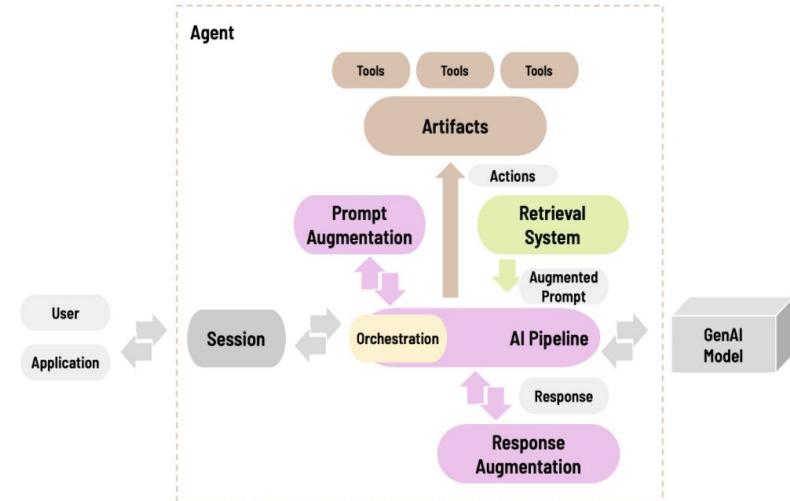
## Dynamic Orchestration of Session Control (2/2)

- **Multi-Session Linking** – If a user revisits a previous topic in a new session, then recall relevant knowledge across sessions.
- **User Preference Adaptation** – If the session history shows tone, detail, or style preferences, then apply them to future responses.
- **Conflict Resolution** – If contradictions exist in prior session context, then reconcile or clarify before prompting.
- **Confidence-Based Use** – If confidence in the relevance of session memory is low, then exclude it from the prompt.



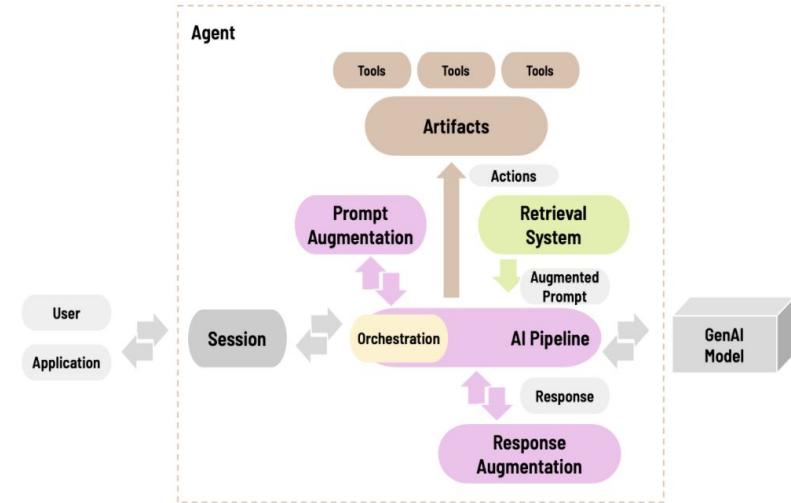
# Dynamic Orchestration of Prompt Augmentation (1/2)

- **Template Switching** – If the task type is Q&A, summarization, reasoning, (other) then switch to the matching prompt template.
- **Role Conditioning** – If the task requires a specific persona, then add system role instructions (e.g., tutor, programmer).
- **Task Decomposition** – If the prompt is too complex, then break it into multiple chained sub-prompts.



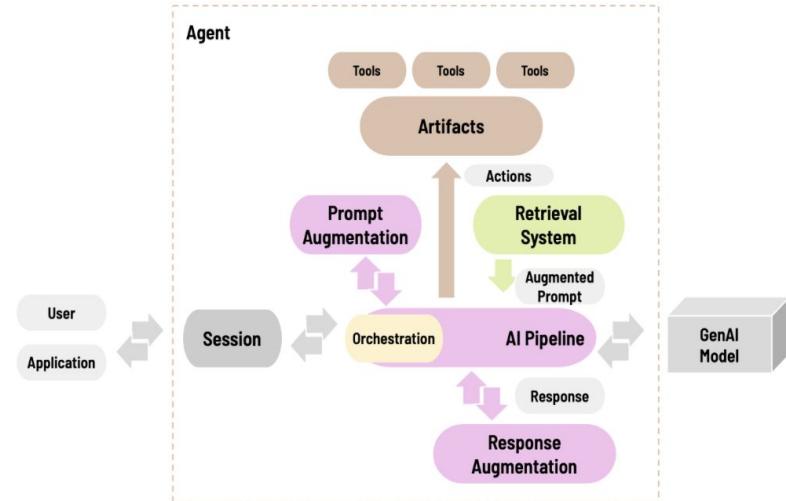
## Dynamic Orchestration of Prompt Augmentation (2/2)

- **Style Adaptation** – If the session history shows a preference for tone, detail, or format, then adapt the style accordingly.
- **Multi-Perspective Framing** – If the task benefits from diverse reasoning, then reframe the prompt in multiple ways before calling the model.
- **Confidence Injection** – If the query is ambiguous or requires accuracy, then strengthen the prompt with validation or reasoning-check instructions.



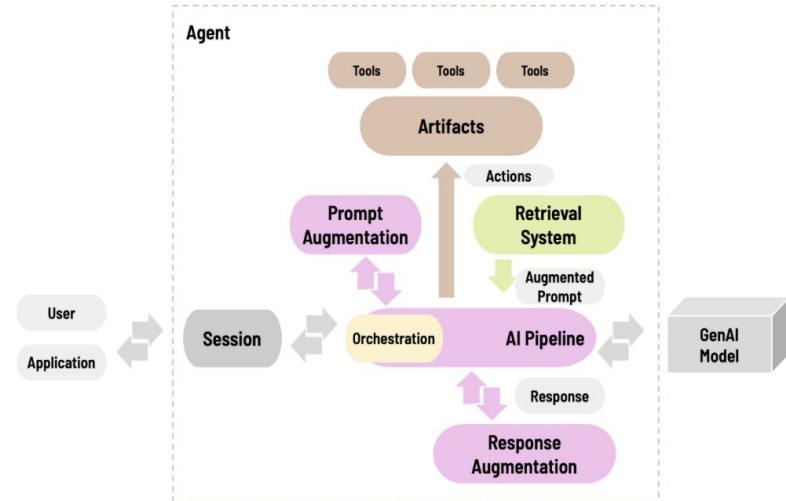
# Dynamic Orchestration of Retrieval Systems (1/2)

- **Adaptive Retrieval Depth** – If the task is complex, then fetch more documents; if it is simple, then fetch fewer.
- **Source Switching** – If the query requires prior context, then use session memory; if it needs external facts, then use a knowledge base or APIs.
- **Hybrid Search** – If the query is factual, then prioritize keyword search; if it is conceptual, then prioritize semantic (vector) search.
- **Multi-Hop Retrieval** – If one retrieval step is insufficient, then perform chained retrieval across multiple sources.
- 



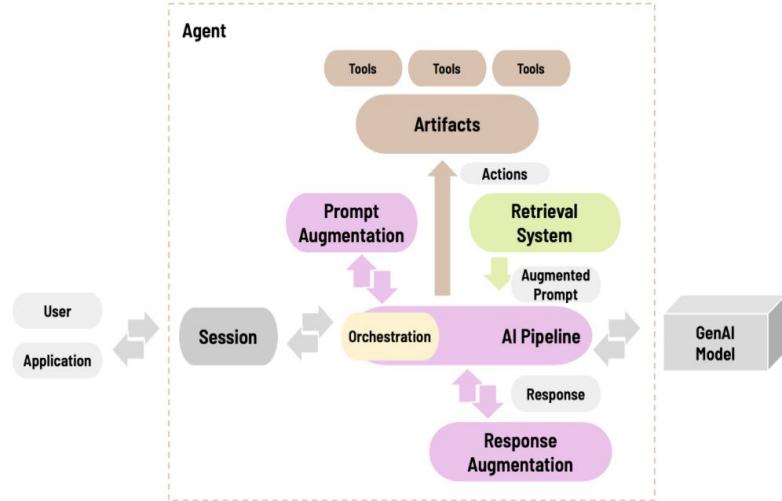
## Dynamic Orchestration of Retrieval Systems (2/2)

- **Relevance Thresholding** – If precision is critical, then apply strict filters for highly relevant documents; if coverage is needed, then relax filters.
- **Domain-Aware Routing** – If the query is domain-specific (e.g., legal, medical, code), then direct it to the appropriate specialized source.
- **Freshness Prioritization** – If the task is time-sensitive, then prioritize newer sources over older ones.
- **Fallback Mechanism** – If the initial retrieval results are poor, then retry with alternative retrieval strategies.



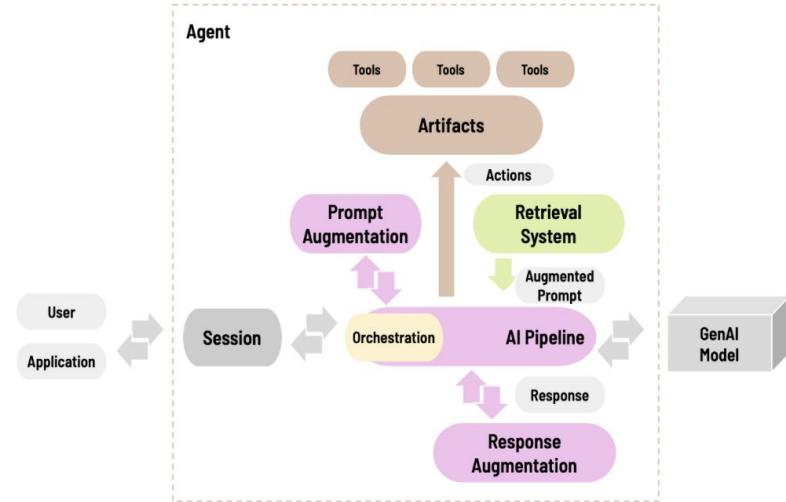
# Dynamic Orchestration of Model Parameters (1/3)

- **Temperature Control** – If the task requires precision, then lower the temperature; if it requires creativity, then raise it.
- **Max Tokens Scaling** – If the task requires detailed output, then increase max tokens; if it requires brevity, then reduce them.
- **Decoding Strategy Switching** – If accuracy is needed, then use greedy decoding; if completeness is needed, then use beam search; if variety is needed, then use sampling.



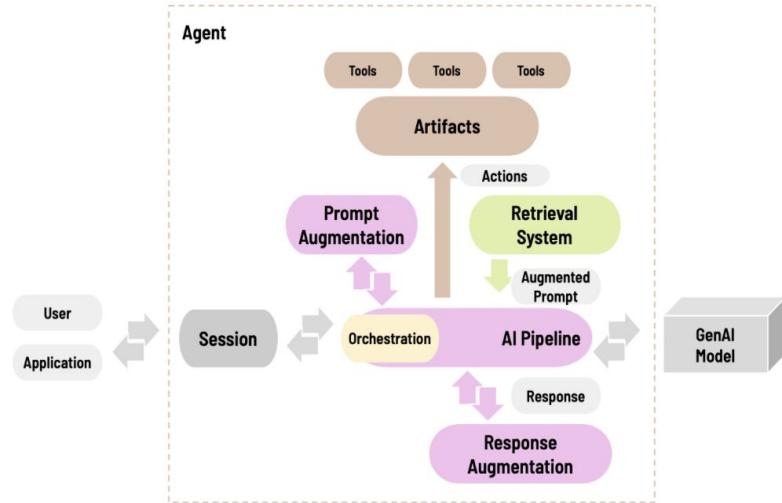
# Dynamic Orchestration of Model Parameters (2/3)

- **Top-k / Top-p Tuning** – If the task requires controlled diversity, then adjust top-k/top-p values accordingly.
- **Detail Level Adaptation** – If the user prefers summaries, then limit verbosity; if they prefer depth, then expand explanations.
- **Confidence-Aware Parameters** – If confidence is low (e.g., weak retrieval scores or uncertainty signals), then increase reasoning depth (more tokens, add self-check/validation steps); if confidence is high, then keep default concise settings.



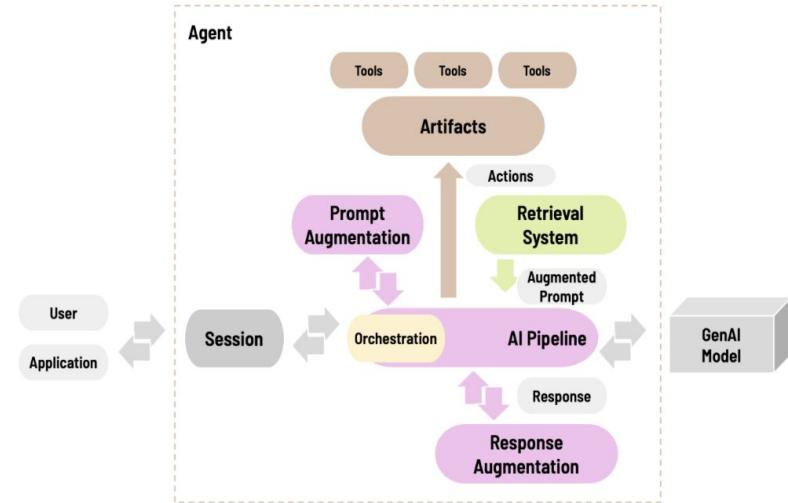
# Dynamic Orchestration of Model Parameters (3/3)

- **Domain-Specific Settings** – If the task is technical/compliance-heavy, then lower temperature and enforce structured output; if it's creative, then raise temperature and allow freer sampling.
- **Progressive Refinement** – If generating from scratch, then draft with looser settings and refine with stricter ones; if polishing existing text, then start strict and iterate lightly.



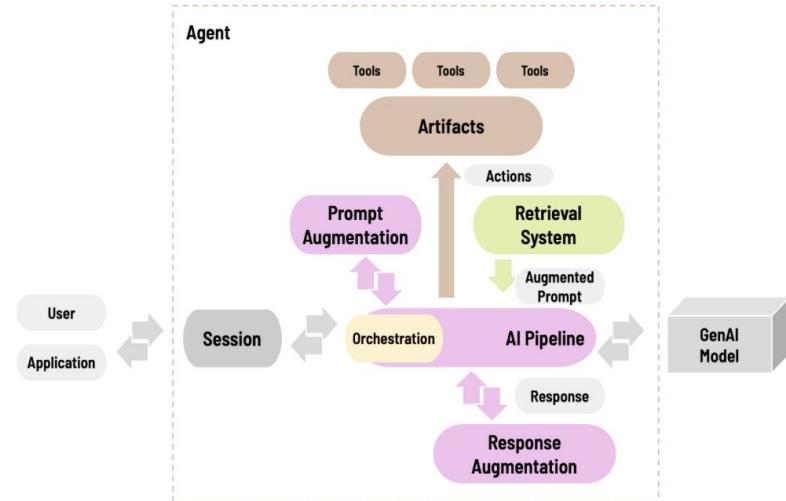
# Dynamic Orchestration of Model Selection (1/2)

- **Lightweight Routing** – If the task is simple and factual, then use a smaller, faster model.
- **Complex Reasoning** – If the task involves multi-step reasoning, then route to a larger, more capable model.
- **Domain-Specific Choice** – If the query is in a specialized domain (e.g., legal, medical, coding), then select a domain-specific model.



## Dynamic Orchestration of Model Selection (2/2)

- **Multi-Model Validation** – If accuracy is critical, then generate answers from multiple models and compare/aggregate.
- **Cost-Aware Selection** – If budget is constrained, then prioritize cheaper models unless high accuracy is required.



---

# Adaptive Orchestration



Can we automate the configuration of GenAI Pipelines?

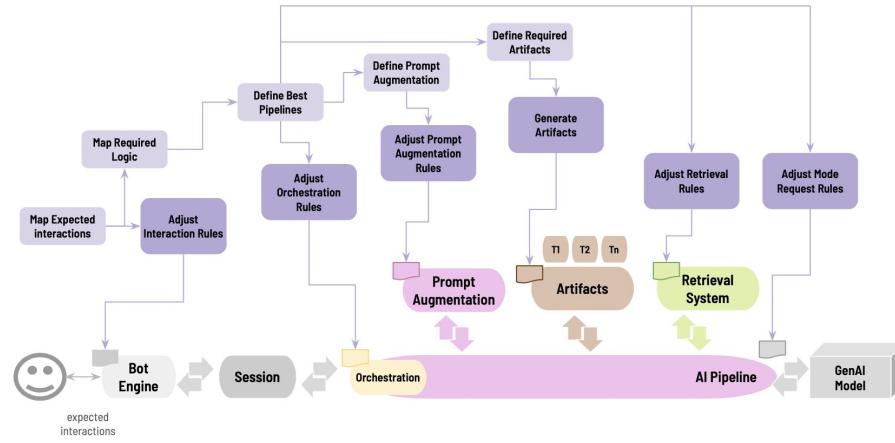
*Would automated configuration be a meaningful step towards advancing AI Solutions?*



# Why to Adapt AI Pipelines?

Adaptation makes AI pipelines more flexible, efficient, and personalized by dynamically adjusting to context and user needs.

- **Self-Programming:** pipelines automatically refine their own rules and strategies
- **Less Coding:** Reduces reliance on hard-coded logic
- **Flexibility:** Adjusts to new tasks and contexts on-the-fly
- **Efficiency:** Reduces wasted steps, latency, and resource costs.



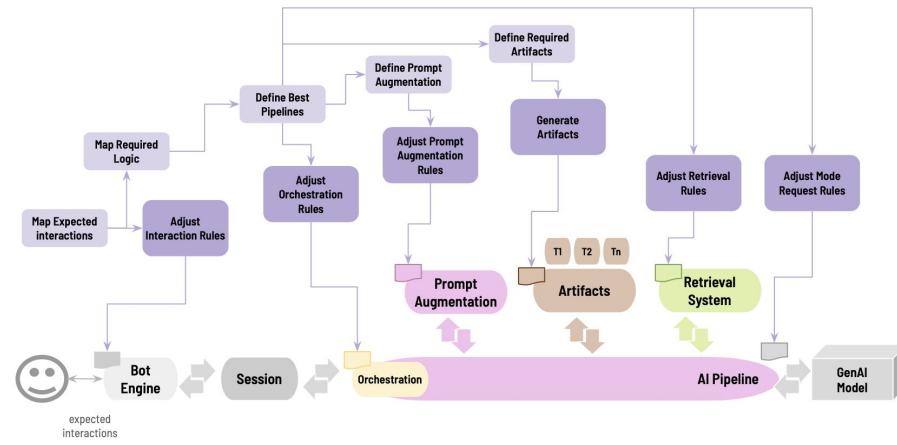
# Why to Adapt in AI Pipelines?

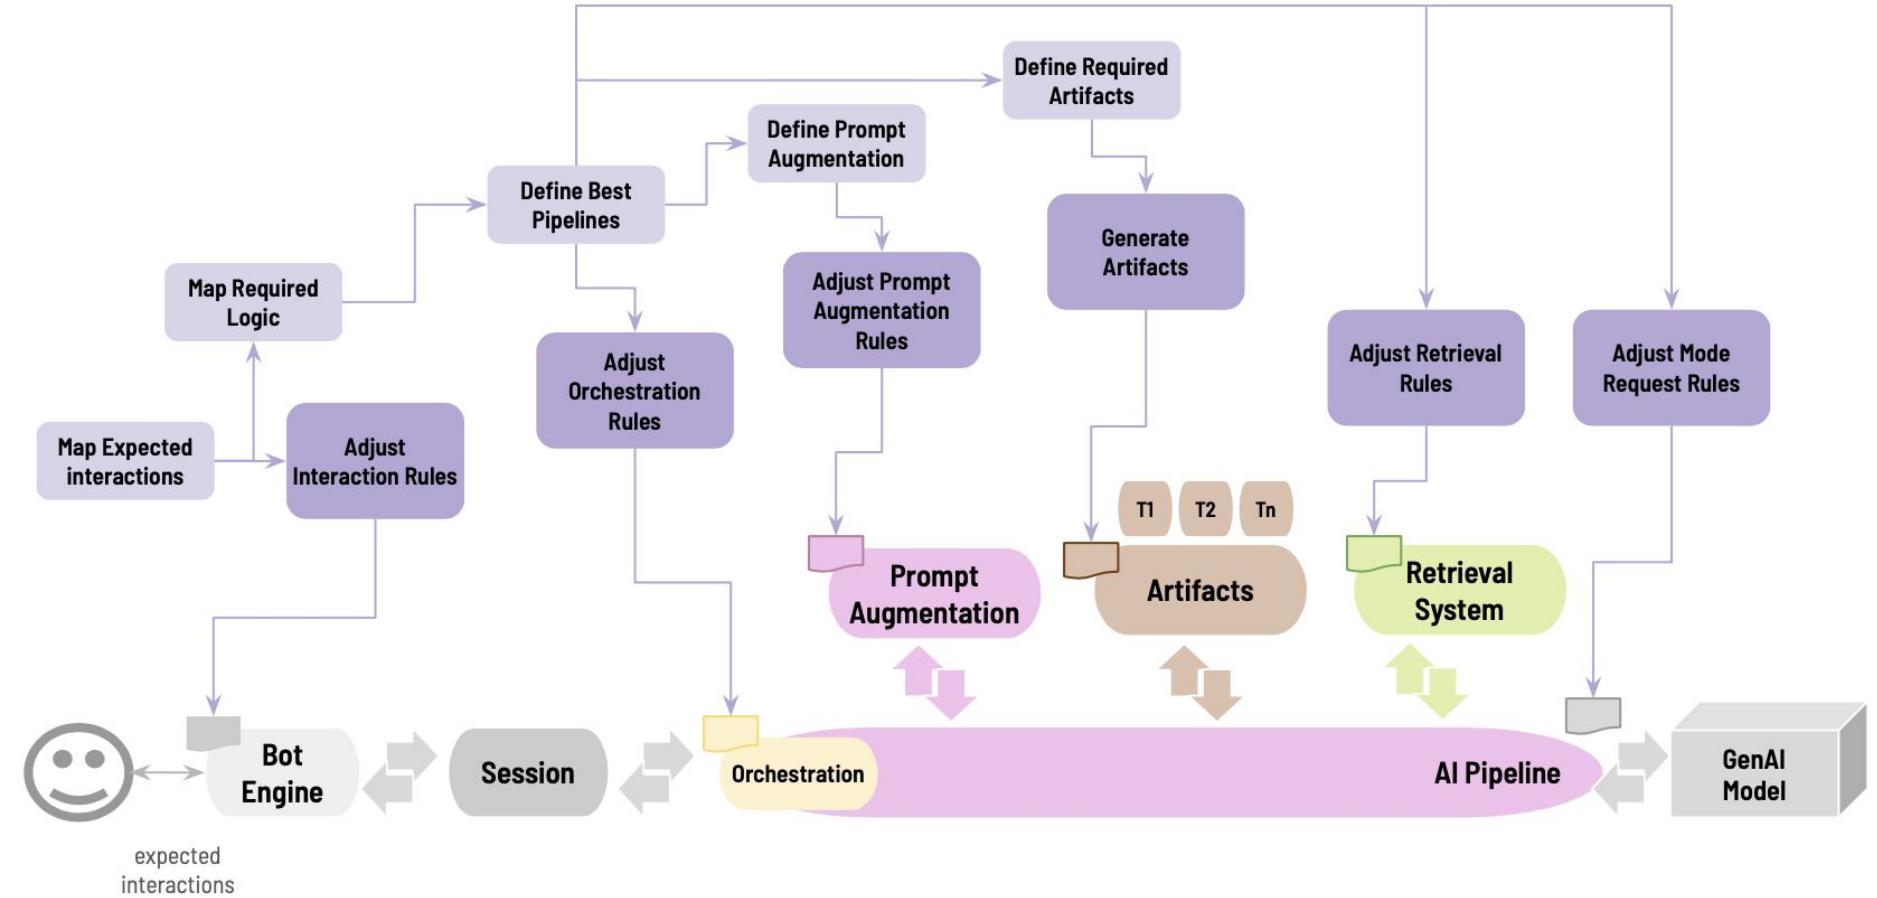
**Adapt Prompt Augmentation Rules** → How can prompts be automatically adapted to improve clarity, personalization, and reasoning?

**Adapt Retrieval Rules** → How can retrieval strategies adapt on-the-fly to provide the right knowledge at the right time?

**Adapt Model Request Rules** → How can the system create new model request strategies for emerging needs?

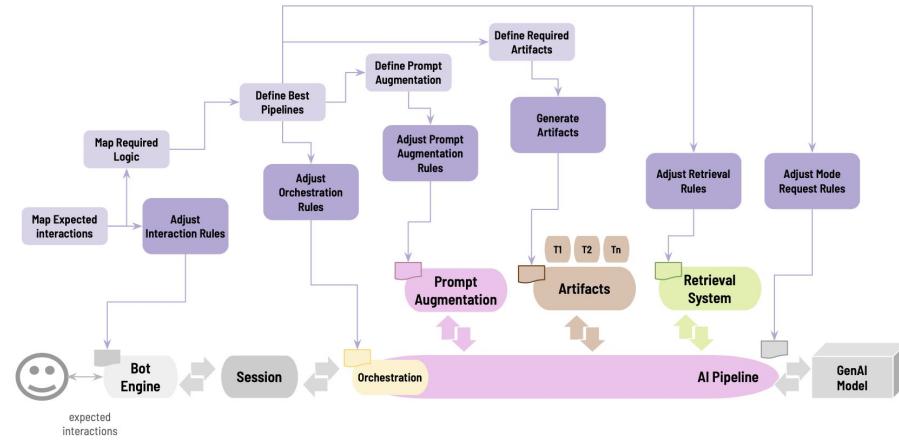
**Generate Artifacts** → How can the system generate new forms of artifacts to improve reasoning and continuity?





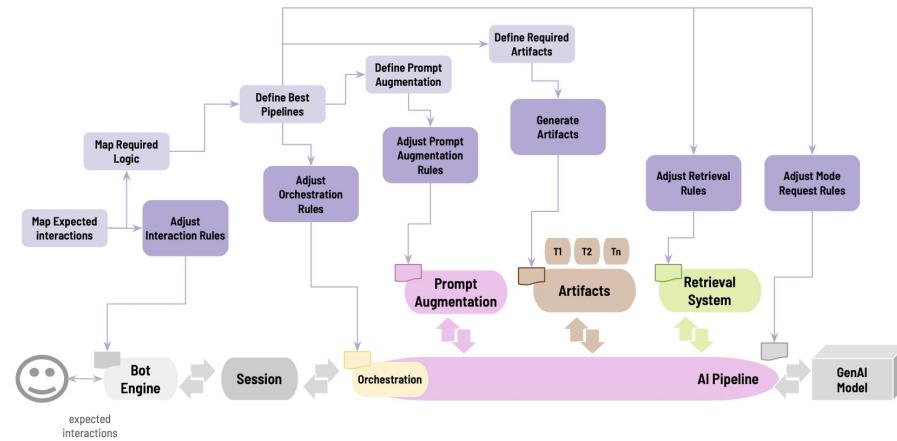
# Adapt Interaction Rules (1/2)

- **Adjust Rule-Based Resolution** – If an interaction can be expressed as condition–action logic, then automatically generate the rules to answer it without AI Pipeline.
- **Optimization Enforcement** – If a rule-based path consistently resolves a query, then strengthen operational rules to bypass AI Pipeline.
- **Preference Adaptation** – If user history shows consistent style/format needs, then build rules that automatically enforce those preferences.



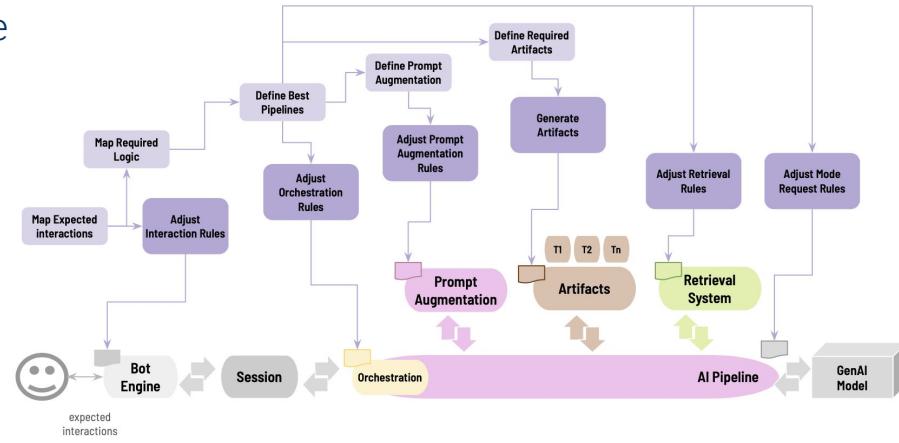
## Adapt Interaction Rules (2/2)

- **Generate API/DB Mappings** – If a user request maps to structured data, then auto-create condition–action rules to call APIs or databases.
- **Optimize Workflow Triggers** – If a query aligns with a known process (e.g., ticket creation), then generate or update rules to directly trigger workflows.
- **Clarification Injection** – If ambiguity is detected in input, then produce a rule that inserts clarification prompts before escalating to AI.
- **Sentiment Conditioning** – If sentiment analysis flags frustration, then adjust or create rules for empathetic, pre-defined responses.



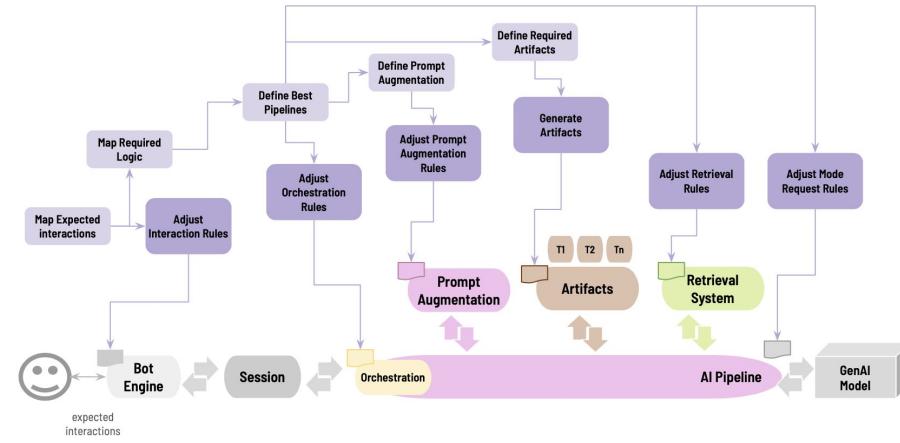
# Adapt Orchestration Rules (1/2)

- **Path Prioritization** – If multiple orchestration paths can satisfy a task, then generate rules to prioritize the most efficient one.
- **Task Decomposition** – If a task requires multi-step reasoning, then create rules to chain prompts dynamically.
- **Intent Shift** – If user intent changes mid-task, then adapt rules to reroute execution.
- **Consensus Building** – If multiple outputs are available, then generate rules for validation and selection.



## Adapt Orchestration Rules (2/2)

- **Load Balancing** – If workload spikes, then create rules to distribute across models/tools.
- **Latency Optimization** – If speed is critical, then create rules to skip optional steps.
- **Conflict Resolution** – If orchestration rules clash, then enforce precedence rules automatically.



---

# EXERCISE



---

## Exercise 5 - Getting Started with MyBot

**Objective:** Implement a modular AI pipeline, orchestrating data inputs, model inferences, and post-processing stages into a coherent flow.

**[Go to Canvas -> Assignments](#)**



# COT 6930 - Generative Intelligence and Software Development Lifecycles

**Dr. Fernando Koch**

kochf@fau.edu

<http://www.fernandokoch.me>